

INF553 Foundations and Applications of Data Mining

Fall 2019

Competition Project

Submission 1 Deadline: November 18th 11:59 PM PST

Submission 2 Deadline: December 2nd 11:59 PM PST

Final Submission Deadline: December 13th 11:59 PM PST

1. Overview of the Competition Project

In this competition project, you need to improve the performance of your recommendation system in Assignment 3. The dataset you are going to use is still the subsets from the Yelp dataset (<https://www.yelp.com/dataset>) used in Assignments 3. You can use any method (like the hybrid recommendation systems) to improve the prediction accuracy and efficiency. There are mandatory bi-weekly submissions on November 18th & December 2nd. The final project submission is on December 13th.

2. Competition Requirements

2.1 Programming Language and Library Requirements

- a. **You must use Python to implement the competition project.** You can use external Python libraries (e.g., Numpy, Pandas, or Scikit-learn).
- b. You are required to only use the Spark RDD for Spark operations. You will not receive any point if you use Spark DataFrame or DataSet.

2.2 Programming Environment

We will use **Python 3.6 and Spark 2.3.3** to test your code. There will be no point if we cannot run your code due to the library version inconsistency.

2.3 Write your own code

Do not share your code with other students!!

We will combine all the code we can find from the Web (e.g., GitHub) as well as other students' code from this and other (previous) sections for plagiarism detection. We will report all the detected plagiarism.

2.4 What you need to turn in

Your submission must be a **zip file** with the naming convention: **firstname_lastname_competition.zip** (all lowercase, e.g., tommy_trojan_competition.zip). You should pack the following required (and optional) files in a folder named **firstname_lastname_competition** (all lowercase, e.g., tommy_trojan_competition) in the zip file (Figure 1, only the files in the red boxes are required to submit):

- a. **[REQUIRED]** Python scripts containing the main function, named:
firstname_lastname_competition.py
- b. **[REQUIRED]** A description file that describes the method you are using (less than 300 words), named:
firstname_lastname_description.txt
- c. **[OPTIONAL]** You can include other scripts to support your programs (e.g., callable functions), but you need to make sure after unzipping, they are all in the same folder “firstname_lastname_competition”.
- d. You don’t need to include any result. We will grade your code using our testing data. Our testing data will be in the same format as the validation dataset.

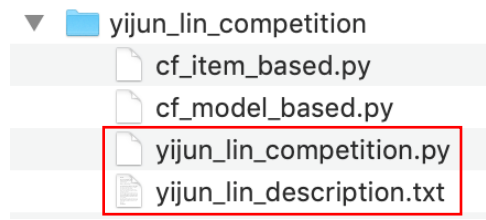


Figure 1: The folder structure after your submission file is unzipped.

3. Yelp Data

In this competition, the datasets you are going to use are in the Google drive:

<https://tinyurl.com/y5vqq6j4>

We have generated the datasets A, B, and C from the original Yelp review dataset with some filters such as the condition: “state” == “CA”. We randomly took 60% of the data as the training dataset, 20% of the data as the validation dataset, and 20% of the data as the testing dataset.

A. yelp_train.csv: the training data, which only include the columns: user_id, business_id, and stars.

B. yelp_val.csv: the validation data, which are in the same format as the training data.

C. We do not share the testing dataset.

D. Other datasets: providing additional information (like the location of a business)

- a. review_train.json: review data only for the training pairs (user, business)
- b. user.json: all user metadata
- c. business.json: all business metadata, including locations, attributes, and categories
- d. checkin.json: user checkins for individual businesses
- e. tip.json: tips (short reviews) written by a user about a business
- f. photo.json: photo data, including captions and classifications

4. Task (8 points)

In the competition, you need to build a recommendation system to predict the given (user, business) pairs. You can mine interesting and useful information from the datasets provided in the Google Drive folder “Competition” to support your recommendation system.

You must make an improvement to your recommendation system in terms of **accuracy**. You can utilize the validation dataset (yelp_val.csv) to evaluate the accuracy of your recommendation system. There are two options to evaluate your recommendation system:

1) Error Distribution: You can compare your results with the corresponding ground truth and compute the absolute differences. You can divide the absolute differences into 5 levels and count the number for each level as the following:

>=0 and <1: 12345
>=1 and <2: 123
>=2 and <3: 1234
>=3 and <4: 1234
>=4: 12

This means that there are 12345 predictions with < 1 difference from the ground truth. This way you will be able to know the error distribution of your predictions and to improve the performance of your recommendation system.

2) RMSE Error: You can compute the RMSE (Root Mean Squared Error) by using following formula:

$$RMSE = \sqrt{\frac{1}{n} \sum_i (Pred_i - Rate_i)^2}$$

Where $Pred_i$ is the prediction for business i and $Rate_i$ is the true rating for business i . n is the total number of the business you are predicting.

Input format: (we will use the following command to execute your code)

```
Python: $ ./bin/spark-submit firstname_lastname_competition.py <input_path> <test_file_name> <result_file_name>
```

Param <input_path>: the path of the data folder (e.g., Competition/), which contains the exact same files on the Google drive

Param <test_file_name>: the name of the testing file (e.g., yelp_val.csv), including the file path

Param: <result_file_name>: the name of the prediction result file, including the file path

Note: Do not hardcode the file paths inside your code.

Output format:

a. The output file is a CSV file, containing **all the prediction results for each user and business pair** in the validation/testing data. The header is “user_id, business_id, prediction”. There is no requirement for the order in this task. There is no requirement for the number of decimals for the similarity values. Please refer to the format in Figure 2.

```

user_id, business_id, prediction
C5QsUsQg5I3dMdLM02SXGA,PvGyzCh1PTga4ePE2-iB2Q,5.0
oxd0FmY0YWW4gFq5jJr-hg,ZSCEkqlzZKRrZUz98CXtNw,2.804287677476818
GGTF7hnQi6D5W77_qiKlqg,5PyqkF8zZbfgFDyAcLUehQ,4.688318401935079

```

Figure 2: Output example in CSV

b. You also need to submit a **txt file** includes the description of your method (less than 300 words) with the naming convention “firstname_lastname_description.txt” (e.g., tommy_trojan_description.txt). The description file should include the explanation to the models you are using, especially the way you improve the accuracy or efficiency of the system. We look forward to seeing creative methods. **Please also report the error distribution, RMSE, and the total execution time on the validation dataset in the description. Figure 3 shows an example of the description file.**

```

Method Description:
...

Error Distribution:
>=0 and <1: 12345
>=1 and <2: 123
>=2 and <3: 1234
>=3 and <4: 1234
>=4: 12

RMSE:
1.11

Execution Time:
200s

```

Figure 3: An example of description file

Grading:

We will compare your prediction results against the ground truth. we will use **our testing data** to evaluate your recommendation systems and grade based on the accuracy using RMSE.

To get the full points for the competition project, **your RMSE result should beat TAs’**. TAs will also continuously improve their systems and will announce the RMSE baseline for the validation data every Friday. They will fix their final results **on December 7th**. However, if your recommendation system only beats the TAs’ for the validation data, you will receive **50% of the points for the competition**.

You should post your accuracy result for the validation data on the Discussion board “Competition project” and compete with TAs’ and other students’ results. The final submission with the highest accuracy will receive extra **3 points on the final grade**. The second place will receive extra **2 points**. The third one will receive extra **1 point**.

All three submissions contribute towards the total project competition score. Grading Distribution for the three submissions are as follows:

Submission 1 (November 18th): 25%

Submission 2 (December 2nd): 25%

Final Submission (December 13th): 50%

NOTICE: Current RMSE baseline is 1.12 for the validation dataset.

Note: RMSE baseline for Submission 2 will be posted on November 19th.

5. Grading Criteria

(% penalty = % penalty of possible points you get)

1. You cannot use the extension for the competition. No late for the competition.
2. If we cannot run your programs with the command we specified (including version conflicts and library issues), you will not receive point for the competition.
3. If the header of the output file is missing, there will be 20% penalty.
4. We will not regrade on the competition.
5. There will be no point if the total execution time exceeds **20 minutes**.