

Finding Similar Sets

Applications

Shingling

Minhashing

Locality-Sensitive Hashing

New thread: High dim. data

High dim. data

Locality sensitive hashing

Clustering

Dimensionality reduction

Graph data

PageRank, SimRank

Network Analysis

Spam Detection

Infinite data

Filtering data streams

Web advertising

Queries on streams

Machine learning

SVM

Decision Trees

Perceptron, kNN

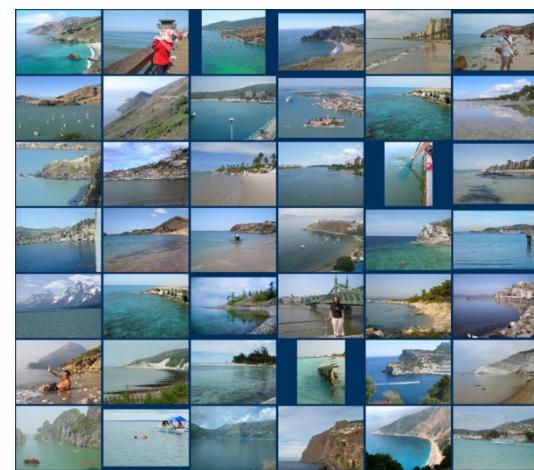
Apps

Recommender systems

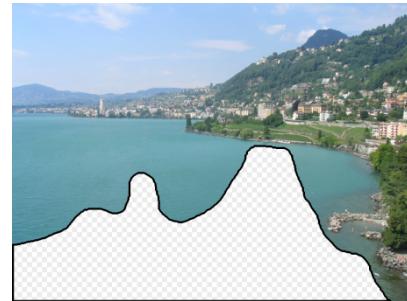
Association Rules

Duplicate document detection

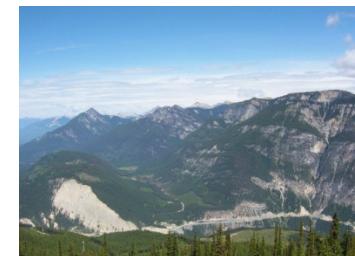
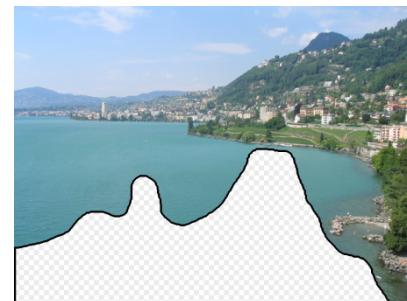
Scene Completion Problem



Scene Completion Problem



Scene Completion Problem



10 nearest neighbors from a collection of 20,000 images

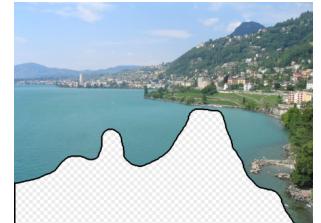
Scene Completion Problem



10 nearest neighbors from a collection of 2 million images⁶

A Common Metaphor

- ◆ Many problems can be expressed as finding “similar” sets:
 - Find near-neighbors in high-dimensional space
- ◆ Examples:
 - Pages with similar words
 - For duplicate detection, classification by topic
 - Movie Rating, NetFlix users with similar tastes in movies
 - For recommendation systems
 - Customers who purchased similar products
 - Products with similar customer sets
 - Images with similar features
 - Users who visited similar websites.



Problem for Today's Lecture

◆ Given: High dimensional data points x_1, x_2, \dots

➤ For example: Image is a long vector of pixel colors

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 2 & 1 \\ 0 & 1 & 0 \end{bmatrix} \rightarrow [1\ 2\ 1\ 0\ 2\ 1\ 0\ 1\ 0]$$

◆ And some distance function $d(x_1, x_2)$

➤ Which quantifies the “distance” between x_1 and x_2

◆ Goal: Find all pairs of data points (x_i, x_j) that are within some distance threshold $d(x_i, x_j) \leq s$

◆ Note: Naïve solution would take $O(N^2)$ ☹

➤ where N is the number of data points

➤ $O(N^2)$ represents an algorithm whose performance is directly proportional to the square of the size of the input data set

◆ MAGIC: This can be done in $O(N)$!! How?

Finding Similar Items

Finding Similar Documents

- ◆ Given a body of **documents**, e.g., the Web, find **pairs of documents** with a lot of **text in common**, such as:
 - **Mirror sites**, or approximate mirrors,
 - **Application:** Don't want to show both in a search.
 - **Plagiarism**, including large quotations
 - **Similar news articles** at many news sites,
 - **Application:** Cluster articles by “same story.”

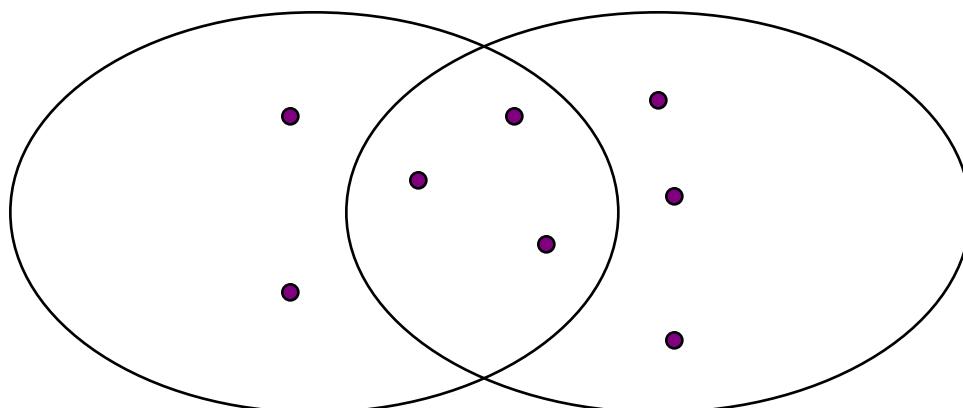
Distance Measures

- **Goal: Find near-neighbors in high-dimensional space**
 - We formally define “near neighbors” as points that are a “small distance” apart
- ◆ For each application, we first need to define what “distance” means
- ◆ **Today: Jaccard distance/similarity.**

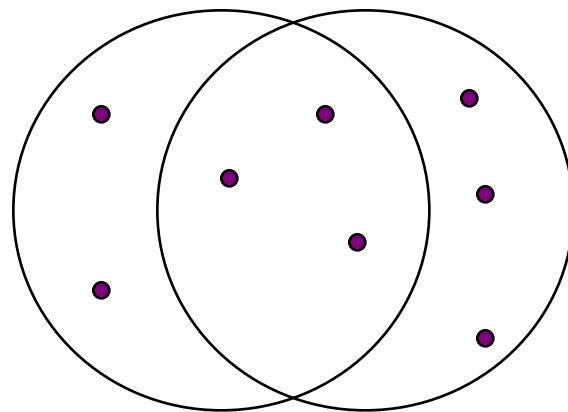
Jaccard Similarity of Sets

- ◆ The *Jaccard similarity* of two sets is the size of their intersection divided by the size of their union

$$Sim(C_1, C_2) = |C_1 \cap C_2| / |C_1 \cup C_2|$$



Example: Jaccard Similarity



3 in intersection
8 in union

Jaccard similarity = 3/8

- **Jaccard distance** = $1 - \text{Jaccard Similarity}$ or
 $5/8$ in this example

Application: Collaborative Filtering

- ◆ Online purchase
 - Recommend items from other buyers
 - ⇒ Need to find similar buyers or items
 - ⇒ Buyer-to-item or item-to-item recommendation
- ◆ Buyer = a set of items he/she purchased
 - Similar buyers: users who bought many common items
- ◆ Item = a set of buyers who purchased it
 - Similar items: items having many common buyers.

Collaborative Filtering @Amazon

◆ Item-to-item recommendation

Customers Who Bought This Item Also Bought

Page 1 of 19



Photive iPad Air Smart Case. Lightweight Smart Cover Case for the New iPad Air with Built in...

★★★★★ 848

\$19.95



iPad Air Case, SUPCASE Heavy Duty Beetle Defense Series Full-body Rugged Hybrid Protective Case...

★★★★★ 2,193

\$19.99



Tech Armor Apple iPad Air 2 / iPad Air (first generation) High Definition (HD) Clear Screen...

★★★★★ 4,686

#1 Best Seller in Tablet Screen Protectors
\$9.95



Tech Armor Apple iPad Air 2 / iPad Air (first generation) High Definition (HD) Clear Screen...

★★★★★ 2,471

\$9.99



Fintie iPad Air Case - SmartShell Case for Apple iPad Air (iPad 5) 2013 Model, Ultra Slim...

★★★★★ 1,795

\$14.99



简 中文

Finding Similar Buyers

Who is most similar to B1?

B1	{A, B}
B2	{A, B, D}
B3	{A, C, D}

$$\text{Jaccard}(B1, B2) = 2/3$$
$$\text{Jaccard}(B1, B3) = 1/4$$

Buyer-to-item recommendation

B2 most similar to B1
B2 also bought D

Recommend **D** to **B1**

Who is most similar to B2?
 $\text{Jaccard}(B2, B1) = \text{Jaccard}(B1, B2)$
 $\text{Jaccard}(B2, B3) = 1/2$

Who is most similar to B3?

$$\text{Jaccard}(B3, B1) = \text{Jaccard}(B1, B3)$$
$$\text{Jaccard}(B3, B2) = \text{Jaccard}(B2, B3)$$

Finding Similar Items

B1	{A, B}
B2	{A, B, D}
B3	{A, C, D}

Buyer as a set



A	{B1, B2, B3}
B	{B1, B2}
C	{B3}
D	{B2, B3}

Item as a set

Finding Similar Items

Which item is most similar to A?

A	{B1, B2, B3}
B	{B1, B2}
C	{B3}
D	{B2, B3}

Which item is most similar to B?

Finding Similar Items

Which item is most similar to A?

A	{B1, B2, B3}
B	{B1, B2}
C	{B3}
D	{B2, B3}

$$\begin{aligned}\text{Jaccard}(A, \textcolor{red}{B}) &= 2/3 & A \text{ is most similar to } \\ \text{Jaccard}(A, C) &= 1/3 & \text{B and D} \\ \text{Jaccard}(A, \textcolor{red}{D}) &= 2/3\end{aligned}$$

Which item is most similar to B?

Finding Similar Items

Which item is most similar to A?

A	{B1, B2, B3}
B	{B1, B2}
C	{B3}
D	{B2, B3}

Which item is most similar to B?

$$\text{Jaccard}(A, \textcolor{red}{B}) = 2/3 \quad A \text{ is most similar to}$$

$$\text{Jaccard}(A, C) = 1/3 \rightarrow \text{B and D}$$

$$\text{Jaccard}(A, \textcolor{red}{D}) = 2/3$$

$$\text{Jaccard}(B, \textcolor{red}{A}) = \text{Jaccard}(A, B) = 2/3$$

$$\text{Jaccard}(B, C) = 0 \rightarrow \text{B is most}$$

$$\text{Jaccard}(B, D) = 1/3 \rightarrow \text{similar to A}$$

Formulated as frequent itemset problem?

◆ Find similar items

- Items bought together by many users
- ⇒ User = transaction
- ⇒ Similar items = frequent item pair

B1	{A, B}
B2	{A, B, D}
B3	{A, C, D}

◆ Find similar users

- Users that bought many common items
- Item = transaction
- => Find frequent user pairs

A	{B1, B2, B3}
B	{B1, B2}
C	{B3}
D	{B2, B3}

Finding Similar Items

Which item is most similar to A?

A	{B1, B2, B3}
B	{B1, B2}
C	{B3}
D	{B2, B3}

Which item is most similar to B?

B1	{A, B}
B2	{A, B, D}
B3	{A, C, D}

$$\text{Jaccard}(A, \textcolor{red}{B}) = 2/3 \quad A \text{ is most similar to}$$

$$\text{Jaccard}(A, C) = 1/3 \rightarrow \text{B and D}$$

$$\text{Jaccard}(A, \textcolor{red}{D}) = 2/3$$

$$\text{Jaccard}(B, \textcolor{red}{A}) = \text{Jaccard}(A, B) = 2/3$$

$$\text{Jaccard}(B, C) = 0 \rightarrow \text{B is most}$$

$$\text{Jaccard}(B, D) = 1/3 \rightarrow \text{similar to A}$$

Item-to-item recommendation

 B1 has bought A, B (or put them in basket)
A is most similar to B and D
B is most similar to A



Recommend: D to B1

Finding Frequent Item Pairs

B1	{A, B}
B2	{A, B, D}
B3	{A, C, D}

$\text{Sup}(A) = 3$ $\text{min_sup} = 2$
 $\text{Sup}(B) = 2$ 
 $\text{Sup}(C) = 1$ prune
 $\text{Sup}(D) = 2$

Apriori algorithm

Finding Frequent Item Pairs

B1	{A, B}
B2	{A, B, D}
B3	{A, C, D}

$\text{Sup}(A) = 3$ $\text{min_sup} = 2$
 $\text{Sup}(B) = 2$ $\text{Sup}(C) = 1$ $\text{Sup}(D) = 2$ $\xrightarrow{\text{prune}}$

$\text{Sup}(A) = 3$ Candidate 2-itemsets
 $\text{Sup}(B) = 2$ $\text{Sup}(D) = 2$ $\xrightarrow{} (A,B)$
 (A,D)
 (B,D)

Apriori algorithm

Finding Frequent Item Pairs

B1	{A, B}
B2	{A, B, D}
B3	{A, C, D}

$\text{Sup}(A, B) = 2$
 $\text{Sup}(A, C) = 1$ $\text{min_sup} = 2$
 $\text{Sup}(A, D) = 2$ \longrightarrow (A, B)
 $\text{Sup}(B, C) = 0$ (A, D)
 $\text{Sup}(B, D) = 1$
 $\text{Sup}(C, D) = 1$

$\text{Sup}(A) = 3$ $\text{min_sup} = 2$
 $\text{Sup}(B) = 2$ \longrightarrow
 $\text{Sup}(C) = 1$ prune $\text{Sup}(A) = 3$ Candidate 2-itemsets
 $\text{Sup}(D) = 2$ $\text{Sup}(B) = 2$ \longrightarrow (A,B)
 $\text{Sup}(D) = 2$ (A,D)
 (B,D)

Apriori algorithm

Finding Frequent User Pairs

A	{B1, B2, B3}
B	{B1, B2}
C	{B3}
D	{B2, B3}

$\text{Sup}(B1) = 2$ $\text{min_sup} = 2 \rightarrow$
 $\text{Sup}(B2) = 3$ prune
 $\text{Sup}(B3) = 3$

Apriori algorithm

Finding Frequent User Pairs

A	{B1, B2, B3}
B	{B1, B2}
C	{B3}
D	{B2, B3}

$\text{Sup(B1)} = 2$ $\text{min_sup} = 2$
 $\text{Sup(B2)} = 3$
 $\text{Sup(B3)} = 3$ prune

Candidate 2-itemsets	
Sup(B1) = 2	(B1,B2)
Sup(B2) = 3	(B1,B3)
Sup(B3) = 3	(B2,B3)

Apriori algorithm

Finding Frequent User Pairs

A	{B1, B2, B3}
B	{B1, B2}
C	{B3}
D	{B2, B3}

$$\text{Sup}(B1, B2) = 2$$

$$\text{Sup}(B1, B3) = 1$$

$$\text{Sup}(B2, B3) = 2$$

$\text{min_sup} = 2$

(B1, B2)
(B2, B3)

$$\begin{aligned}\text{Sup}(B1) &= 2 \quad \text{min_sup} = 2 \\ \text{Sup}(B2) &= 3 \\ \text{Sup}(B3) &= 3\end{aligned}$$

prune

Candidate 2-itemsets

$$\begin{aligned}\text{Sup}(B1) &= 2 \\ \text{Sup}(B2) &= 3 \\ \text{Sup}(B3) &= 3\end{aligned}$$

\longrightarrow

(B1, B2)
(B1, B3)
(B2, B3)

Apriori algorithm

So why new solution?

- ◆ Just use Apriori to find
 - Frequent item pairs => similar items
 - Frequent user pairs => similar users

B1	{A, B}
B2	{A, B, D}
B3	{A, C, D}

A	{B1, B2, B3}
B	{B1, B2}
C	{B3}
D	{B2, B3}

Potential Problems

- ◆ Are buyer B1 and B2 still similar?
- ◆ What would Apriori say?

B1	{A, B}
B2	{A, B, D}
B3	{A, C, D}
B4	{E, F}
B5	{E, F, G}
B6	{X, Y, Z}
B7	{X, W}
B8	{S, T, O}
B9	{S, T}
...	...

Potential Problems

◆ Are buyer B1 and B2 still similar?

◆ What would Apriori say?

- Form item->buyers basket
- A large number of items
- B1 and B2 only appear together in two baskets
- Hence very low support
- But they are still similar.

B1	{A, B}
B2	{A, B, D}
B3	{A, C, D}
B4	{E, F}
B5	{E, F, G}
B6	{X, Y, Z}
B7	{X, W}
B8	{S, T, O}
B9	{S, T}
...	...

Application: Collaborative Filtering

- ◆ Recommend movies
 - Recommend similar movies or
 - Movies from similar users
- ◆ User = a set of movies he/she has watched
- ◆ Movie = a set of users who has watched it

What if we consider ratings too?

- ◆ User A watched movie HP1 (Harry Potter 1)
 - And rated it 4
- ◆ May need different similarity function/solution

⇒ Recommendation systems (Chapter 9)

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

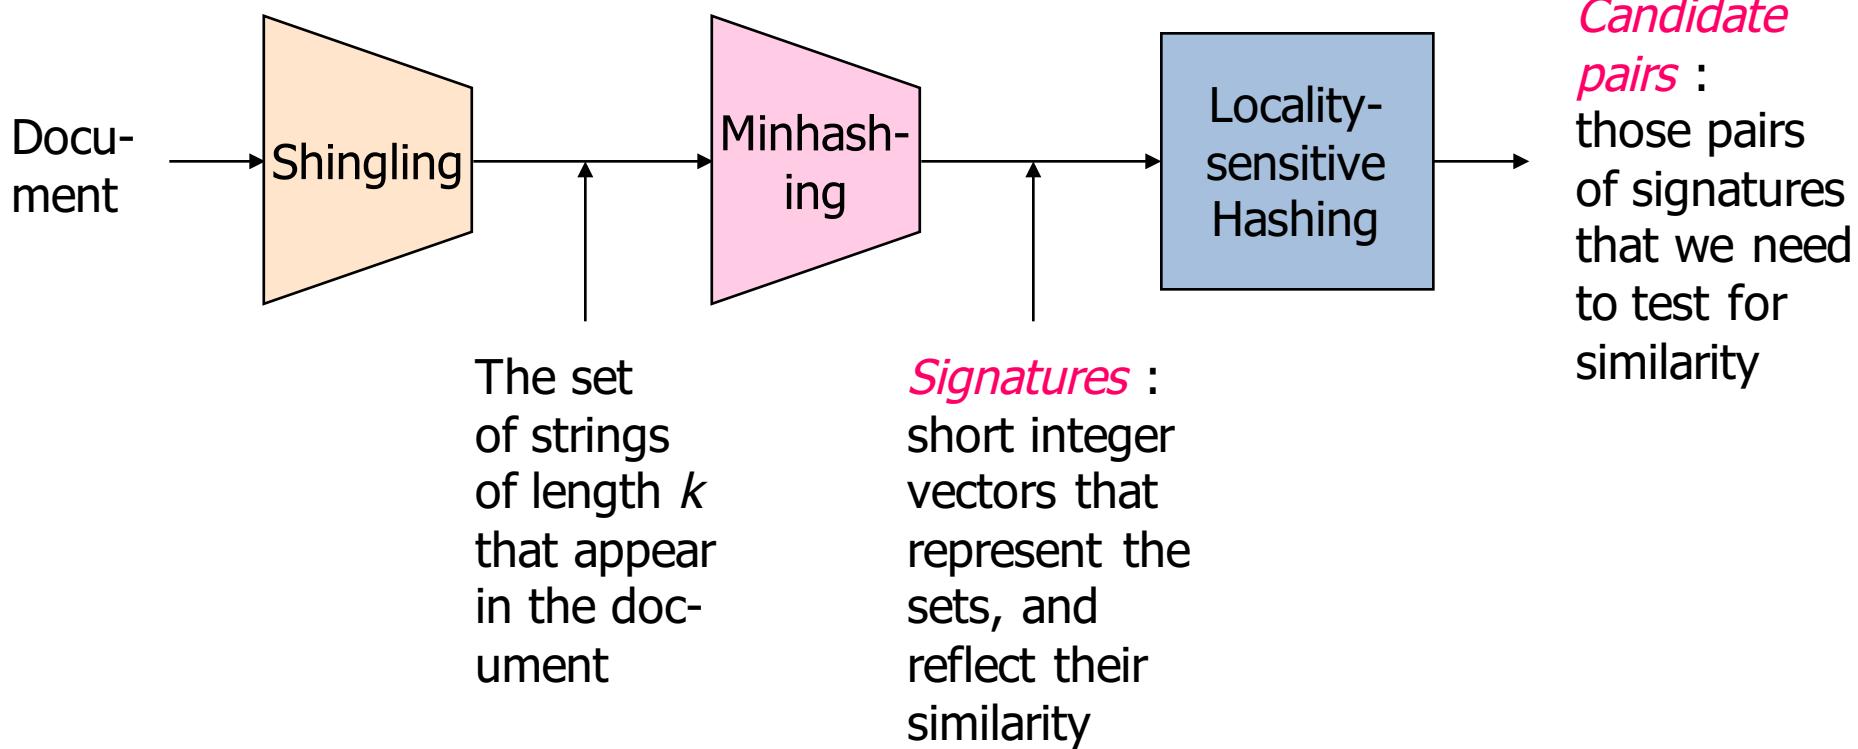
Task: Finding Similar Documents

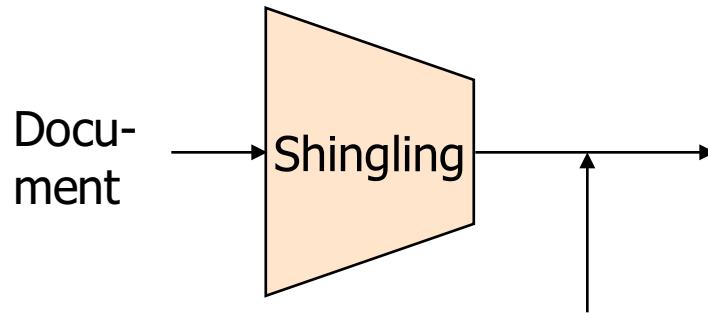
- ◆ **Goal:** Given a large number (in the millions or billions) of documents, find “near duplicate” pairs
- ◆ **Applications:**
 - Mirror websites, or approximate mirrors
 - Don’t want to show both in search results
 - Similar news articles at many news sites
 - Cluster articles by “same story”
- ◆ **Problems:**
 - Many **small pieces** of one document can appear out of order in another
 - **Too many** documents **to compare** all pairs
 - Documents are so large or so many that they **cannot fit in main memory.**

3 Essential Steps for Finding Similar Docs

1. *Shingling*: Convert documents to sets
2. *Min-Hashing*: Convert large sets to short signatures, while preserving similarity
3. *Locality-Sensitive Hashing*: Focus on pairs of signatures likely to be from similar documents
 - Candidate pairs!

The Big Picture





The set
of strings
of length k
that appear
in the doc-
ument

Shingling

Step 1: *Shingling:* Convert documents to sets

Documents as High-Dimensional Data

- ◆ Step 1: *Shingling*: Convert documents to sets
- ◆ Simple approaches:
 - Document = set of words appearing in document
 - Document = set of “important” words
 - Don’t work well for this application. Why?
- ◆ Need to account for ordering of words!
- ◆ A different way: *Shingles*!

Define: Shingles

- ◆ A *k-shingle* (or *k-gram*) for a document is a sequence of k tokens that appears in the doc
 - Tokens can be *characters*, *words* or something else, depending on the application
 - Assume tokens = characters for examples,
- ◆ **Example:** $k=2$; document $D_1 = \text{abcab}$
Set of 2-shingles: $S(D_1) = \{\text{ab}, \text{bc}, \text{ca}\}$
 - **Option:** Shingles as a bag (multiset), count ab twice: $S'(D_1) = \{\text{ab}, \text{bc}, \text{ca}, \text{ab}\}$.

Shingles

- ◆ Web page as a string of characters
- ◆ Shingle = subsequence of k-characters
- ◆ Web page = abcdabd, k = 2
- ◆ 2-shingles
 - ab, bc, cd, da, bd
- ◆ Max # of k-shingles for a page of n characters?
 - 7 characters, 6 max 2-shingles, two identical, 5 unique
 - Max = $7 - 2 + 1 = 6$
 - $N - k + 1 \sim n$

White Spaces

- ◆ Better not omit them
- ◆ Could turn multiple into one
- ◆ D1: “scored a touch down” => “scored a touch down”
- ◆ D2: “touchdown at last”
- ◆ D1 and D2 have a common 9-shingle if space omitted.

Shingle Size

- ◆ Too small
 - Many documents will falsely become similar

- ◆ Too big
 - Might miss truly similar documents.

Working Assumption

- ◆ Documents that have lots of shingles in common have similar text, even if the text appears in different order
- ◆ **Caveat:** You must pick k large enough, or most documents will have most shingles
 - $k = 5$ is OK for short documents (eg. Email, Tweet)
 - $k = 10$ is better for long documents
 - news articles, blog posts (in between)
- ◆ May want to **compress long shingles**.

Compressing Shingles

- ◆ To **compress long shingles**, we can **hash** them to (say) 4 bytes
 - Called ***tokens***
- ◆ **Represent a document by the set of hash values of its *k*-shingles**
 - **Idea:** Two documents could (rarely) appear to have shingles in common, when in fact only the hash-values were shared
- ◆ **Example:** $k=2$; document $D_1 = \text{abcab}$
Set of 2-shingles: $S(D_1) = \{\text{ab}, \text{bc}, \text{ca}\}$
Hash the singles: $h(D_1) = \{1, 5, 7\}$.

Why is compression needed?

◆ How many k-shingles?

- Rule of thumb: imagine **20 characters** in alphabet
 - **26 characters + whitespace = 27**
 - **since "z,q,x" used rarely**
 - A **k-shingle** is a sequence of k consecutive words (or k-gram)
- Estimate of number of k-shingles is 20^k
- 4-shingles: 20^4 or 160,000 or $2^{17.3}$
- 9-shingles: 20^9 or 512,000,000,000 or **2^{39}**

◆ Assume we use **4 bytes** to represent a **bucket**

- ◆ Buckets numbered in range 0 to **$2^{32} - 1$**
- ◆ Much smaller than possible number of **9-shingles** and represent each shingle with 4 bytes, not 9 bytes
 - Compression.

Thought Question

- ◆ Why is it better to hash 9-shingles (say) to 4 bytes than to use 4-shingles?
- ◆ Hint: How random are the 32-bit sequences that result from 4-shingling?

Why hash 9-shingles to 4 bytes rather than use 4-shingles?

- ◆ With 4-shingles, most sequences of four bytes are unlikely or impossible to find in typical documents
- ◆ Effective number of different shingles much less than $2^{32} - 1$
- ◆ With 9-shingles, 2^{39} possible shingles
 - ◆ Many more than 2^{32} buckets
- ◆ After hashing, may get any sequence of 4 bytes.

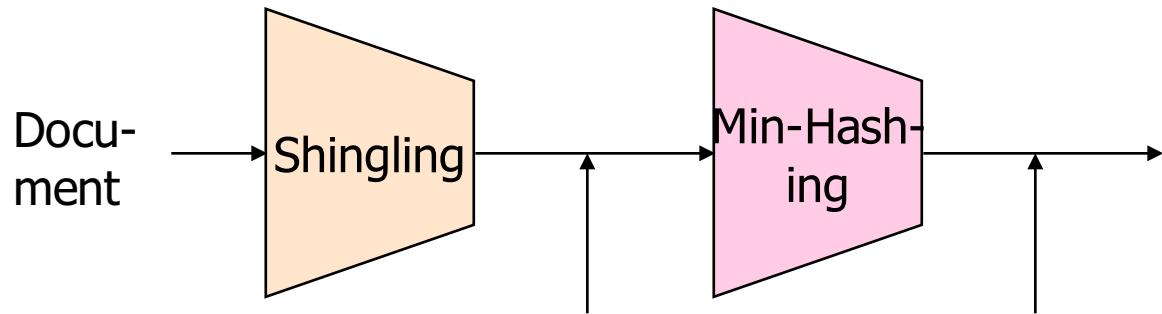
Similarity Metric for Shingles

- ◆ Document D_1 is a set of its k -shingles $C_1 = S(D_1)$
- ◆ Equivalently, each document is a vector of 0s,1s in the space of k -shingles
 - Each unique shingle is a dimension
 - Vectors are very sparse,
- ◆ A natural similarity measure is the Jaccard similarity.

Motivation for Minhash/LSH

Use k-shingles to create Signatures: short integer vectors that represent sets and reflect their similarity

- ◆ Suppose we need to find near-duplicate documents among million documents
- ◆ Naïvely, we would have to compute pairwise Jaccard similarities for every pair of docs
 - For $N = 1$ million
 - i.e, $N(N-1)/2 \approx 5*10^{11}$ comparisons
 - At 10^5 secs/day and 10^6 comparisons/sec, it would take 5 days
- ◆ For $N = 10$ million, it takes more than a year...



The set
of strings
of length k
that appear
in the doc-
ument

Signatures:
short integer
vectors that
represent the
sets, and
reflect their
similarity

MinHashing

Step 2: *Minhashing:* Convert large sets to
short signatures, while preserving similarity

From Sets to Boolean Matrices

- ◆ **Rows** = elements of the universal set
- ◆ **Columns** = sets
 - 1 in **row e** and **column S** if and only if element e is a member of set S
 - Column similarity is the Jaccard similarity of the sets of their rows with 1: intersection/union of sets
- ◆ **Typical matrix is sparse** (many 0 values)
- ◆ May not really represent the data by a boolean matrix
- ◆ Sparse matrices are usually better represented by the list of non-zero values (e.g., triples)
 - But the matrix picture is conceptually useful.

Example 3.6

<i>Element</i>	S_1	S_2	S_3	S_4
<i>a</i>	1	0	0	1
<i>b</i>	0	0	1	0
<i>c</i>	0	1	0	1
<i>d</i>	1	0	1	1
<i>e</i>	0	0	1	0

- ◆ Universal set: {a, b, c, d, e}
- ◆ Matrix represents sets chosen from universal set
- ◆ $S_1 = \{a, d\}$, $S_2 = \{c\}$, $S_3 = \{b, d, e\}$ and $S_4 = \{a, c, d\}$
- ◆ Example: **rows are products and columns are customers**, represented by **set of items** they bought
- ◆ **Jacquard similarity of S_1 , S_4 ?** intersection/union = 2/3

Example: Jaccard Similarity of Columns

C₁ C₂

0 1 *

1 0 *

1 1 * *

$$\text{Sim } (C_1, C_2) = 2/5 = 0.4$$

0 0

1 1 * *

0 1 *

Four Types of Rows

- ◆ Given columns C1 and C2, rows may be classified as:

	C1	C2
a	1	1
b	1	0
c	0	1
d	0	0

- Also, $a = \# \text{ rows of type } a$, etc.
- ◆ Note $\text{Sim}(\text{C1}, \text{C2}) = a/(a + b + c)$.

When Is Similarity Interesting?

1. When the **sets** are **so large** or so many that they **cannot fit in main memory**
2. Or, when there are **so many sets** that **comparing all pairs** of sets takes **too much time**
3. Or both.

Outline: Finding Similar Columns

1. Compute **signatures** of columns = **small summaries** of columns
2. Examine **pairs of signatures** to find similar signatures
 - **Essential:** similarities of signatures and columns are related
3. **Optional:** check that columns with similar signatures are really similar.

Warnings

1. Comparing **all pairs of signatures** may take **too much time**, even if not too much space
 - A job for **Locality-Sensitive Hashing**
2. These methods can produce false negatives, and even false positives (if the optional check is not made).

Signatures

- ◆ Key idea: “hash” each column C to a small *signature* $Sig(C)$, such that:
 1. $Sig(C)$ is **small** enough that we can fit a signature in **main memory** for each column
 2. $Sim(C_1, C_2)$ is the same as the “**similarity**” of $Sig(C_1)$ and $Sig(C_2)$.

Four Types of Rows

- ◆ Given columns C_1 and C_2 , rows may be classified as:

	$\underline{C_1}$	$\underline{C_2}$
a	1	1
b	1	0
c	0	1
d	0	0

- ◆ Also, $a = \# \text{ rows of type } a$, etc.
- ◆ Note $\text{Sim}(C_1, C_2) = a / (a + b + c)$
 - Jacquard similarity: intersection/union
 - a is intersection, $a+b+c$ is union

Minhashing

1. To ***minhash*** a set represented by a column of the matrix, **pick a random permutation of the rows**
2. Define “hash” function $h(C)$ = the number of the first (in the permuted order) row in which column C has 1
3. Use several (e.g., 100) independent hash functions to **create a signature**.

Minhashing Example (3.7)

Element	S_1	S_2	S_3	S_4
a	1	0	0	1
b	0	0	1	0
c	0	1	0	1
d	1	0	1	1
e	0	0	1	0

Permute

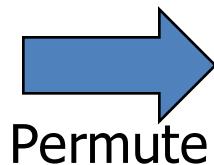
Element	S_1	S_2	S_3	S_4
b	0	0	1	0
e	0	0	1	0
a	1	0	0	1
d	1	0	1	1
c	0	1	0	1

1. To minhash a set represented by a column of the characteristic matrix, pick a **permutation of the rows**
2. The **minhash value** of any column is the number of first row, in permuted order, in which column **has a 1**
3. For set S_1 , first 1 appears in row a , so:

- $h(S_1) =$ 
- $h(S_2) =$ 
- $h(S_3) =$ 
- $h(S_4) =$ 

Minhashing Example (3.7)

<i>Element</i>	S_1	S_2	S_3	S_4
a	1	0	0	1
b	0	0	1	0
c	0	1	0	1
d	1	0	1	1
e	0	0	1	0



<i>Element</i>	S_1	S_2	S_3	S_4
b	0	0	1	0
e	0	0	1	0
a	1	0	0	1
d	1	0	1	1
c	0	1	0	1

- To minhash a set represented by a column of the characteristic matrix, pick a **permutation of the rows**
 - $h(S_1) = a$
 - $h(S_2) = c$
 - $h(S_3) = b$
 - $h(S_4) = a$
- The **minhash value** of any column is the number of first row, in permuted order, in which column **has a 1**
- For set S_1 , first 1 appears in row a , so:

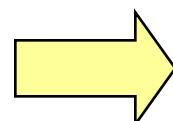
Minhashing Example

Input matrix

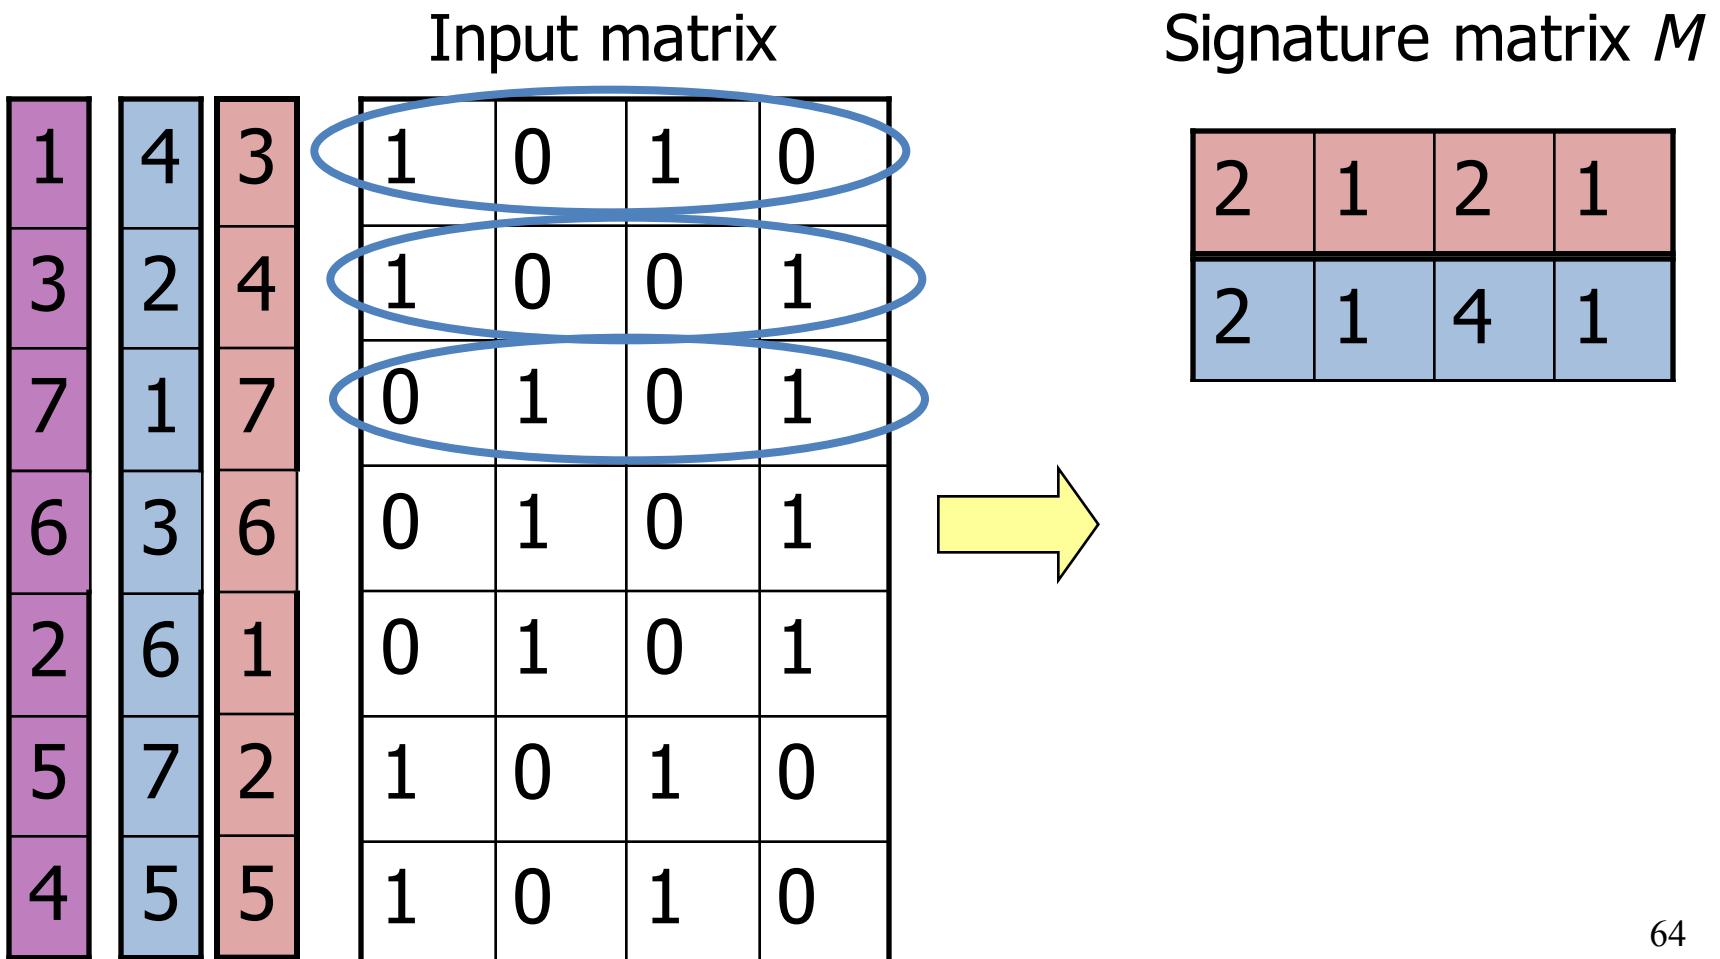
3	1	0	1	0
4	1	0	0	1
7	0	1	0	1
6	0	1	0	1
1	0	1	0	1
2	1	0	1	0
5	1	0	1	0

Signature matrix M

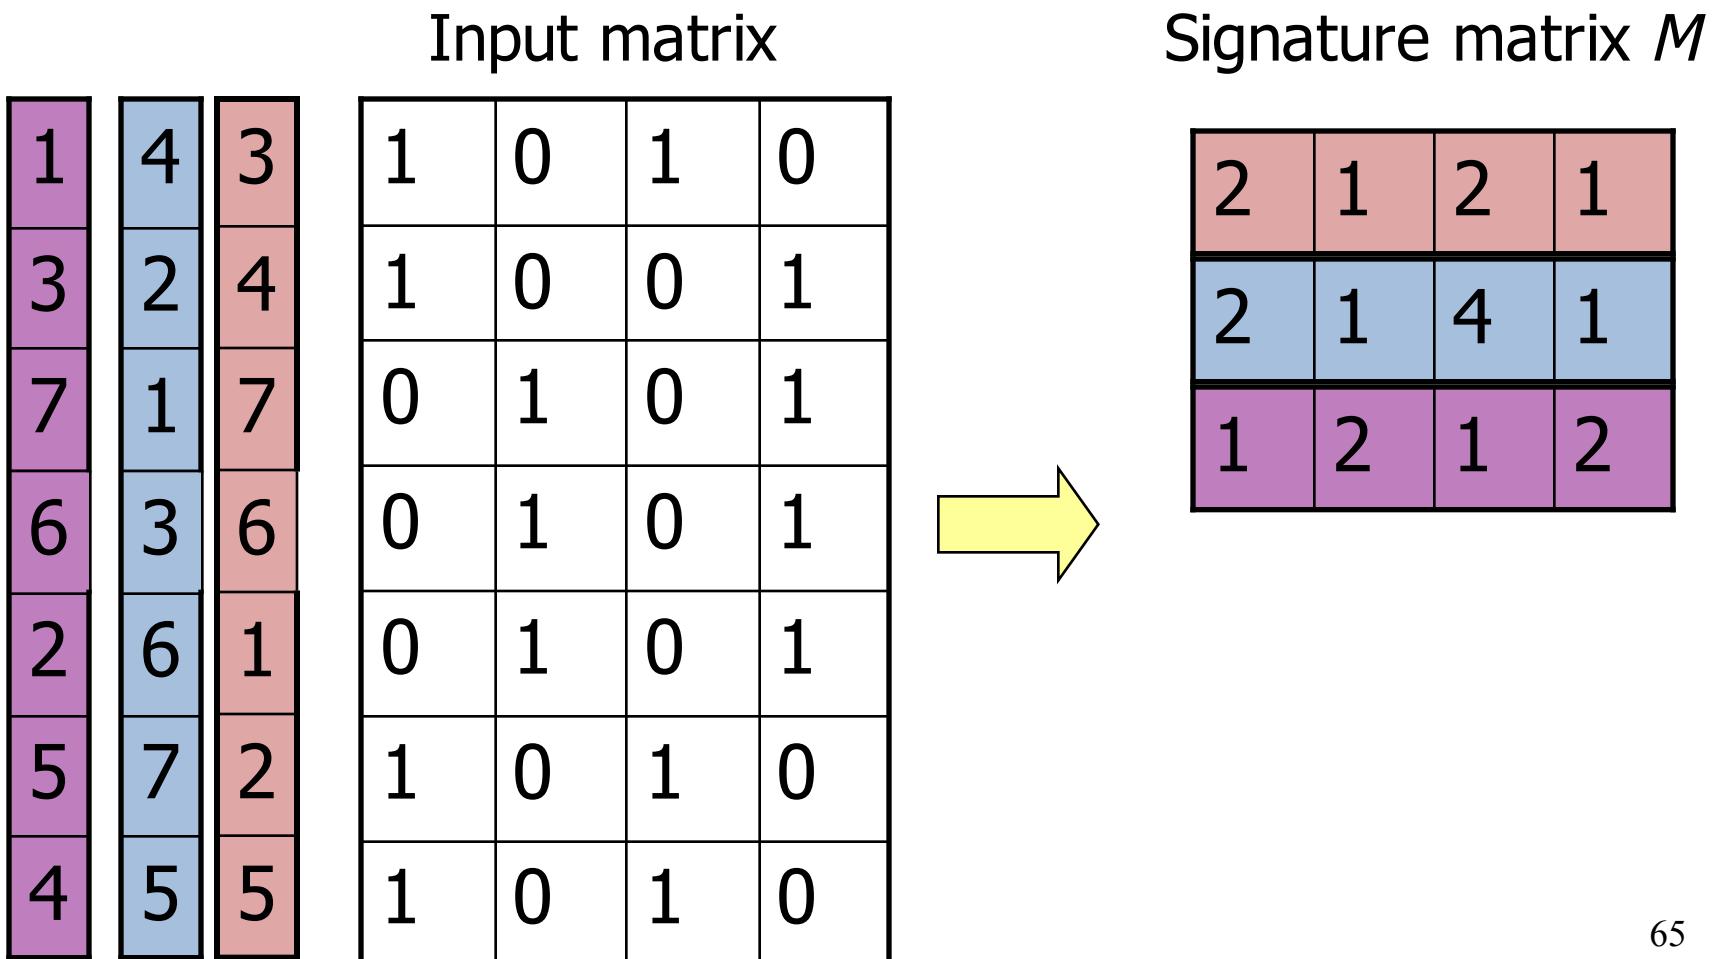
2	1	2	1
---	---	---	---



Minhashing Example



Minhashing Example



Surprising Property: Connection between Minhashing and Jaccard Similarity

- ◆ The probability that minhash function for a **random permutation of rows** produces same value for two sets equals **Jaccard similarity** of those sets
 - Probability that $h(C_1) = h(C_2)$ is the same as $\text{Sim}(C_1, C_2)$
- ◆ Recall four types of rows:

	C_1	C_2
a	1	1
b	1	0
c	0	1
d	0	0

- ◆ $\text{Sim}(C_1, C_2)$ for both Jacquard and Minhash are $a / (a + b + c)$!
 - Why? Look down the permuted columns C_1 and C_2 until we see a 1
 - If it's a type-a row, then $h(C_1) = h(C_2)$. If a type-b or type-c row, then not. (Don't count the type-d rows).

Similarity for Signatures

- ◆ Sets represented by characteristic **matrix M**
- ◆ To represent sets: pick at random some number **n** of permutations of the rows of M
 - 100 permutations or several hundred
- ◆ Call **minhash** functions determined by these permutations h_1, h_2, \dots, h_n
- ◆ From column representing set **S**, construct **minhash signature for S**:
 - vector $[h_1(S), h_2(S), \dots, h_n(S)]$, usually represented as column
- ◆ Construct a ***signature matrix***: **ith column of M replaced by minhash signature for ith column**
- ◆ The ***similarity of signatures*** is the fraction of the hash functions in which they agree.

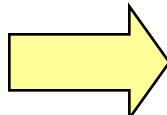
Min Hashing – Example

Input matrix

1	4	3
3	2	4
7	1	7
6	3	6
2	6	1
5	7	2
4	5	5

1	0	1	0
1	0	0	1
0	1	0	1
0	1	0	1
0	1	0	1
1	0	1	0
1	0	1	0

Signature matrix M

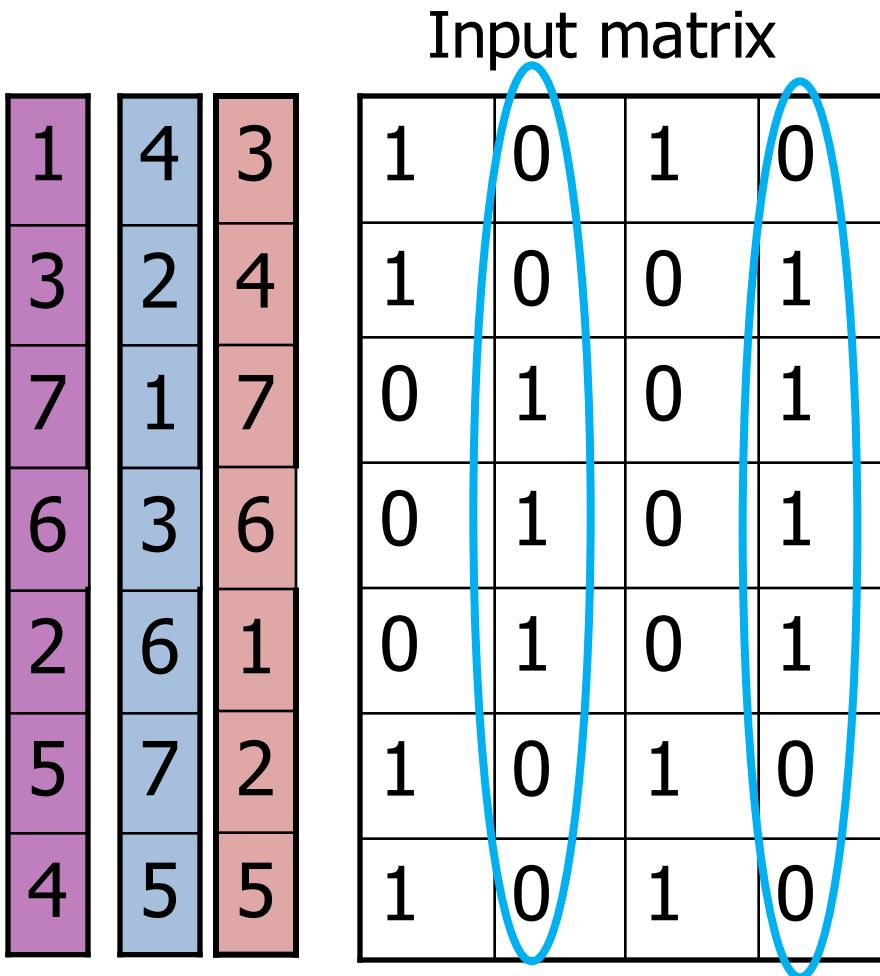


2	1	2	1
2	1	4	1
1	2	1	2

Similarities:

	1-3	2-4	1-2	3-4
Col/Col	0.75	0.75	0	0
Sig/Sig	0.67	1.00	0	0
2/3				

Min Hashing – Example



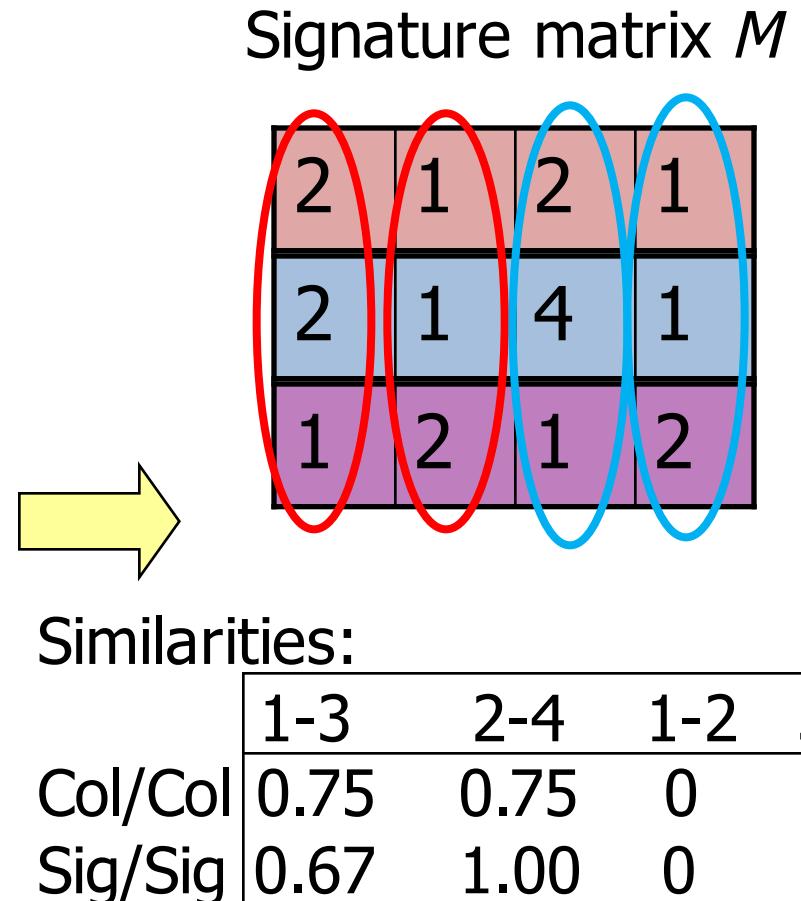
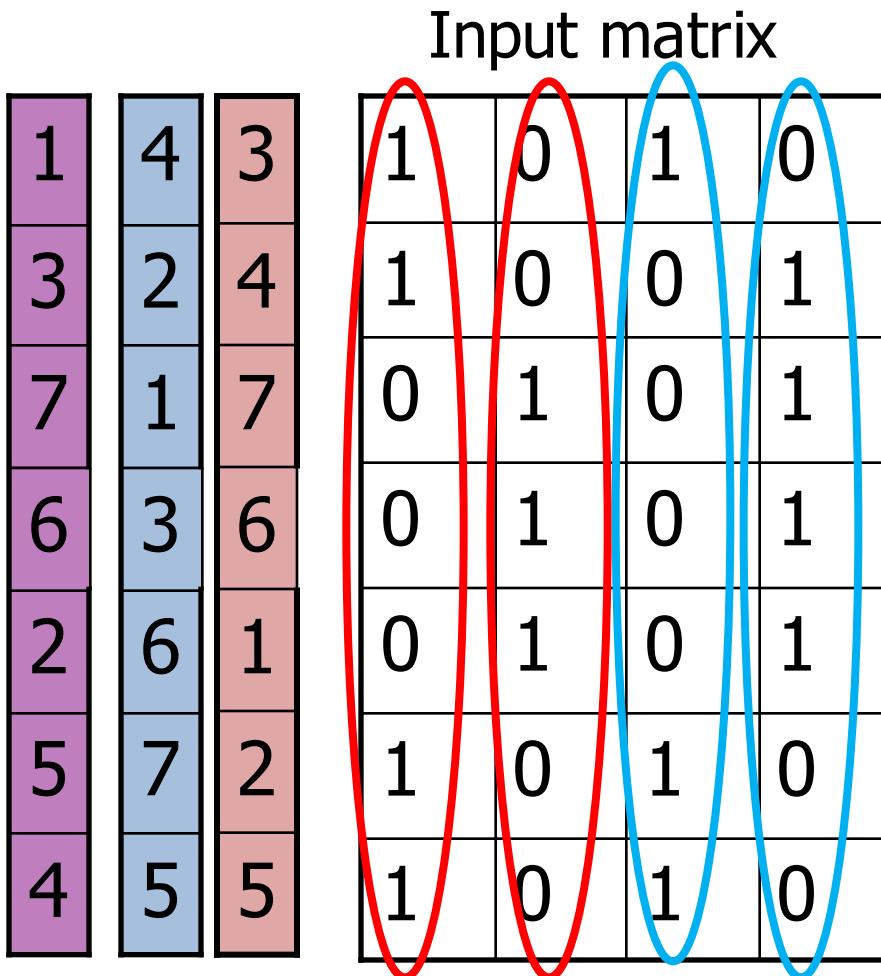
Signature matrix M

2	1	2	1
2	1	4	1
1	2	1	2

Similarities:

	1-3	2-4	1-2	3-4
Col/Col	0.75	0.75	0	0
Sig/Sig	0.67	1.00	0	0

Min Hashing – Example



Minhash Signatures

- ◆ Pick (say) **100 random** permutations of the rows
- ◆ Think of ***Sig* (C)** as a column vector
- ◆ Let ***Sig* (C)[i] =**
according to the ***i* th permutation**, the number of the first row that has a **1 in column C**.

Implementation – (1)

- ◆ **Not feasible to permute** a large characteristic matrix explicitly
 - Suppose **1 billion rows**
 - Hard to pick a **random permutation from 1...billion**
 - Representing a random permutation **requires 1 billion entries**
 - Accessing rows in permuted order leads to thrashing
- ◆ **Can simulate the effect of a random permutation by a random hash function**
 - Maps **row numbers** to as many buckets as there are rows
 - May have **collisions on buckets**
 - **Not important as long as number of buckets is large.**

Implementation – (2)

- ◆ A good approximation to permuting rows:
pick around 100 hash functions
- ◆ For each:
 - column c (set representing a document)
 - hash function h_i
- ◆ Keep a “slot” in signature matrix M (i, c)
- ◆ **Intent:** $M (i, c)$ will become the smallest value of $h_i(r)$ for which column c has 1 in row r
 - $h_i(r)$ gives order of rows for i^{th} permutation.

Implementation – (3)

for each **row r** **do begin**

for each hash function h_i **do**

 Compute $h_i(r)$

for each column c

if c has 1 in row r

for each hash function h_i **do**

if $h_i(r)$ is a smaller value than $M(i, c)$ **then**

$M(i, c) := h_i(r);$

end;

Computing Minhash Signatures:

Example 3.8

Row	S_1	S_2	S_3	S_4	$x + 1 \bmod 5$	$3x + 1 \bmod 5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

Two hash functions give permutations of rows:

$$h_1 = x+1 \bmod 5, h_2 = 3x + 1 \bmod 5$$

	S_1	S_2	S_3	S_4
h_1	∞	∞	∞	∞
h_2	∞	∞	∞	∞

Initial signature matrix

	S_1	S_2	S_3	S_4
h_1	1	∞	∞	1
h_2	1	∞	∞	1

For row 0: Replace existing signature values with lower hash values for S_1 and S_4 , since both have 1 in row 0.

Computing Minhash Signatures:

Example 3.8 (part 2)

Row	S_1	S_2	S_3	S_4	$x + 1 \bmod 5$	$3x + 1 \bmod 5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

	S_1	S_2	S_3	S_4
h_1	1	∞	<u>2</u>	1
h_2	1	∞	<u>4</u>	1

For row 1: replace h_1 and h_2 values for S_3 , since row has a 1 and values are lower

	S_1	S_2	S_3	S_4
h_1	1	<u>3</u>	2	1
h_2	1	<u>2</u>	4	1

For row 2: replace values for S_2 since set has a 1 value. Do not replace values for S_4 , because existing values are lower

Computing Minhash Signatures:

Example 3.8 (part 3)

Row	S_1	S_2	S_3	S_4	$x + 1 \bmod 5$	$3x + 1 \bmod 5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

	S_1	S_2	S_3	S_4
h_1	1	3	2	1
h_2	0	2	0	0

For row 3: don't replace h_1 values--all are below 4; replace h_2 values with 0 for S_1, S_3, S_4

	S_1	S_2	S_3	S_4
h_1	1	3	0	1
h_2	0	2	0	0

For row 4: replace h_1 value for S_3 , don't replace h_2 value since current value is lower

Note: result is same as 77 permutations to find first 1

Example

Row	C1	C2
1	1	0
2	0	1
3	1	1
4	1	0
5	0	1

$$h(x) = x \bmod 5$$

$$g(x) = 2x+1 \bmod 5$$

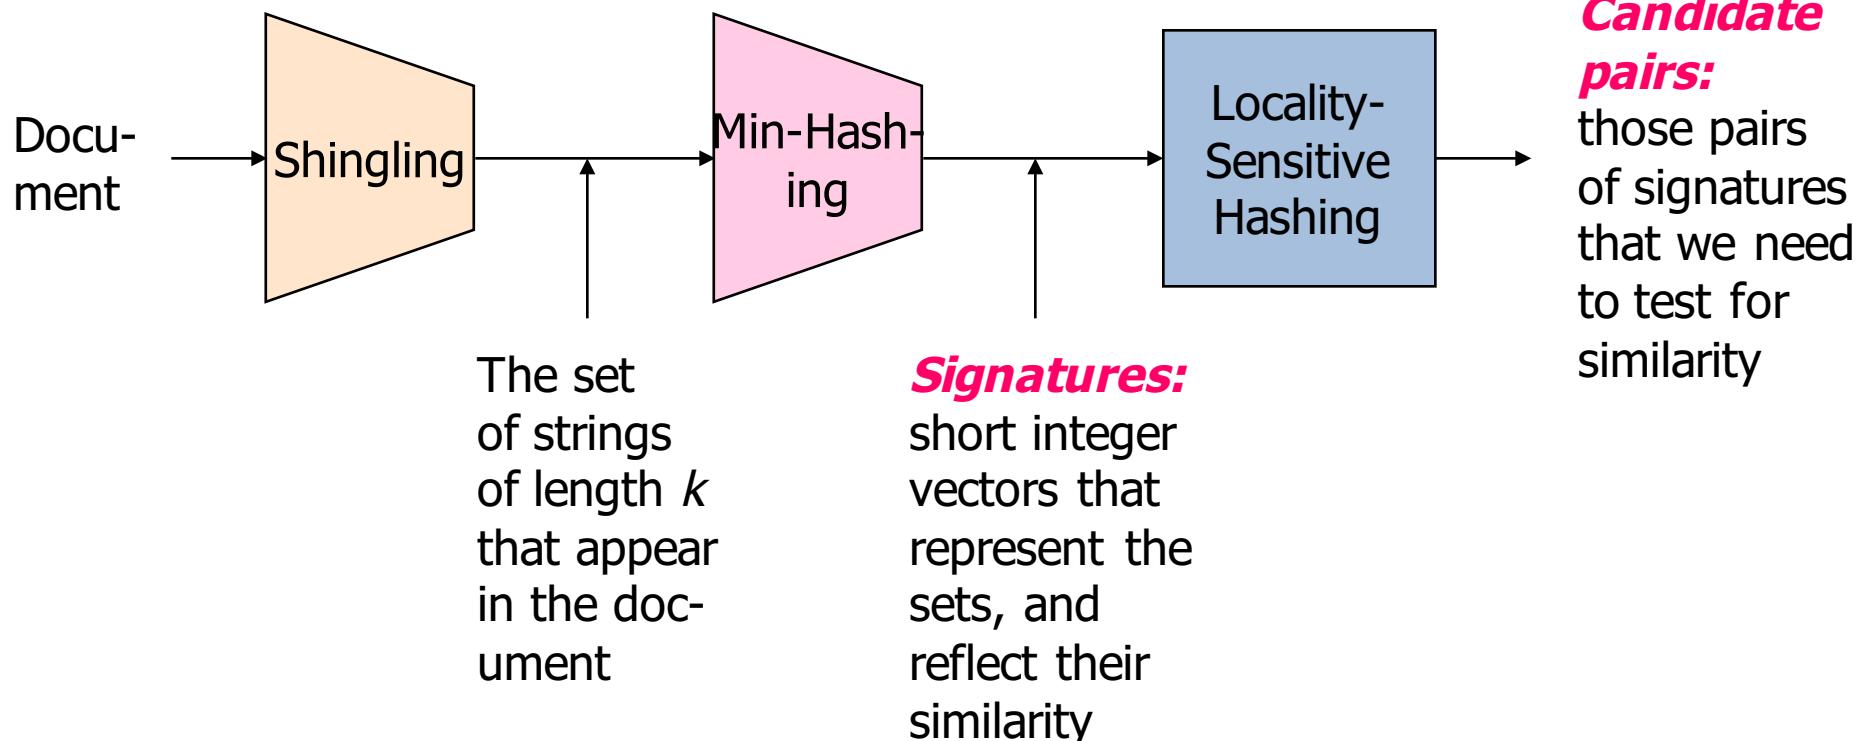
		Sig1	Sig2
	$h(1) = 1$	1	-
	$g(1) = 3$	3	-
	$h(2) = 2$	1	2
	$g(2) = 0$	3	0
	$h(3) = 3$	1	2
	$g(3) = 2$	2	0
	$h(4) = 4$	1	2
	$g(4) = 4$	2	0
	$h(5) = 0$	1	0
	$g(5) = 1$	2	0

Implementation – (4)

- ◆ Often, data is given by column, not row.
 - E.g., columns = documents, rows = shingles.
- ◆ If so, sort matrix once so it is by row.
- ◆ And *always* compute $h_i(r)$ only once for each row.

So far ...

- ◆ Represent a document as a set of hash values (of its k-shingles)
- ◆ Transform set of k-shingles to a set of minhash signatures
- ◆ Use Jaccard to compare two documents by comparing their signatures
- ◆ Is this method (i.e., transforming sets to signature) necessarily “better”??



Locality Sensitive Hashing

Step 3: *Locality-Sensitive Hashing:*

Focus on pairs of signatures likely to be from similar documents