

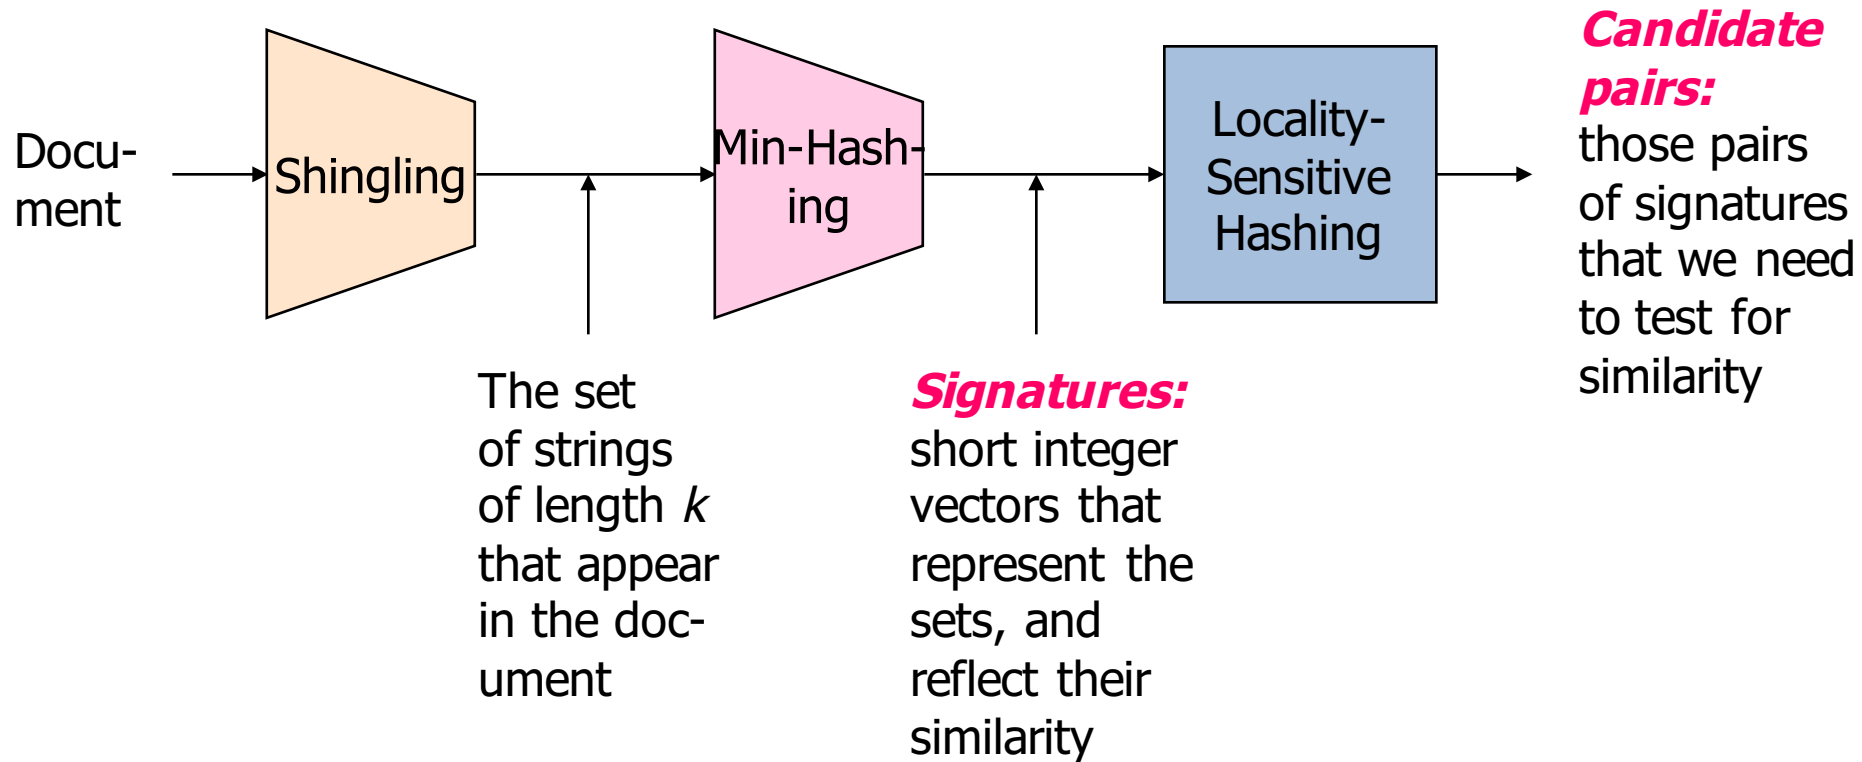
Finding Similar Sets (part 3)

Applications

Shingling

Minhashing

Locality-Sensitive Hashing



Locality Sensitive Hashing

Step 3: **Locality-Sensitive Hashing:**

Focus on **pairs of signatures** likely to be from similar documents

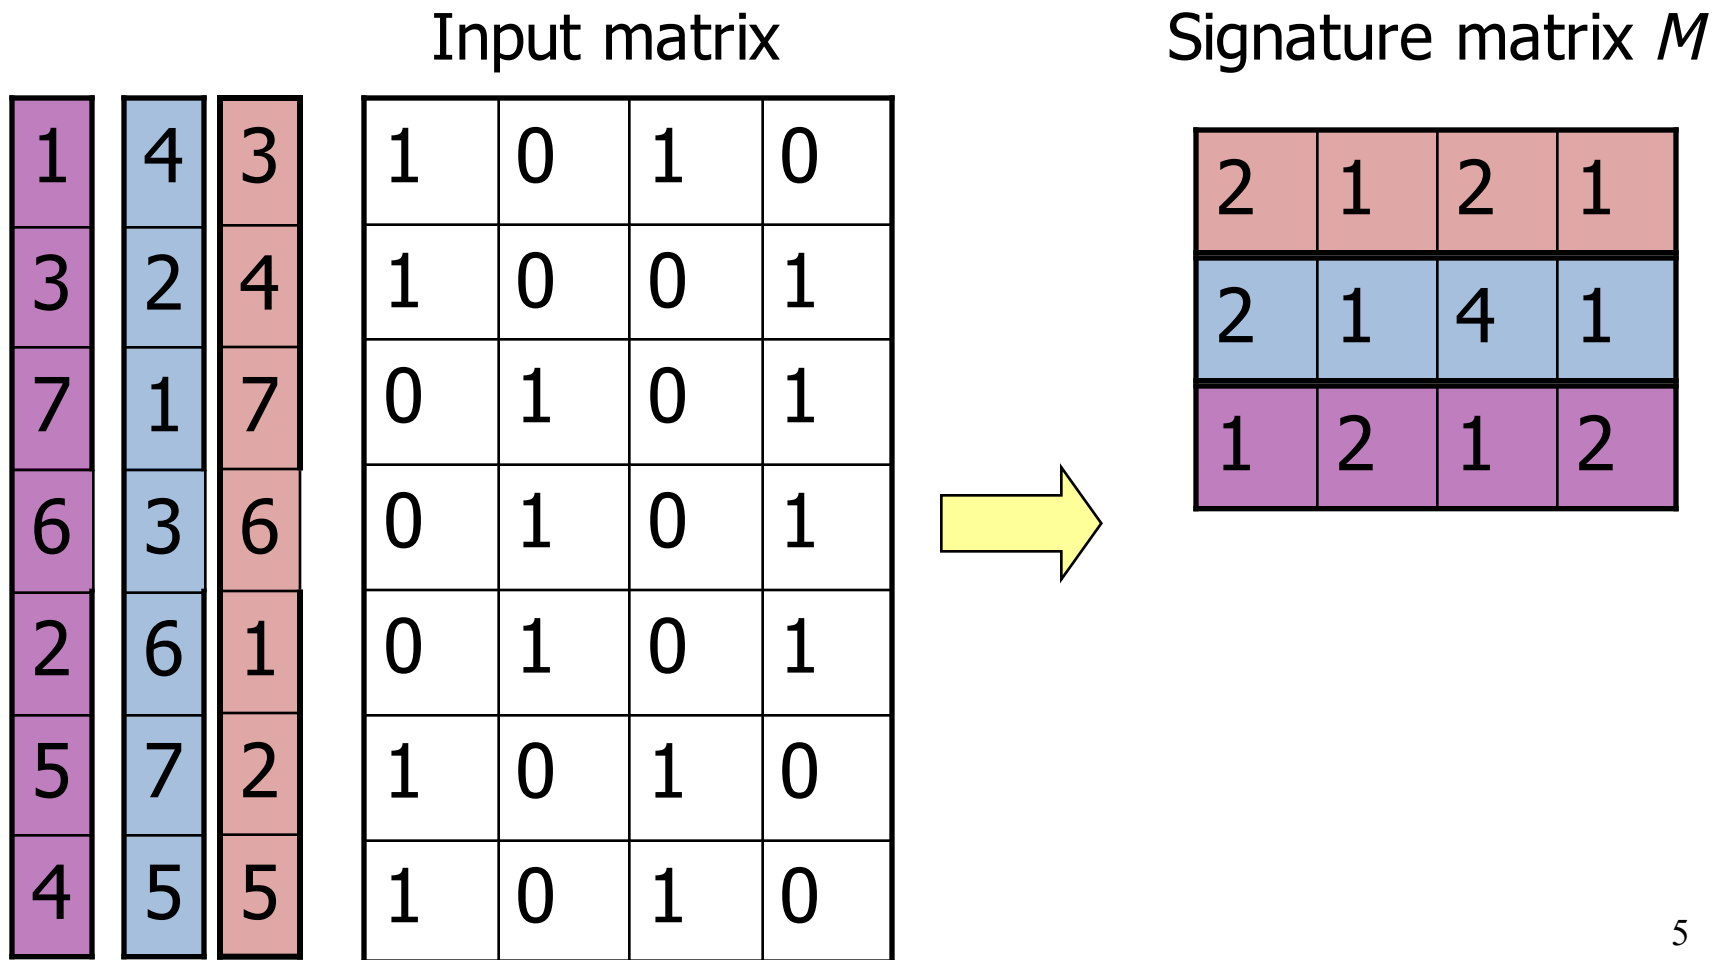
Motivation for Locality Sensitive Hashing

- ◆ *Used **k-shingles** to create sets that summarize documents*
- ◆ *Used **Minhashing** to generate signatures that represent sets of shingles, reflect their similarity*
- ◆ Suppose we need to find near-duplicate documents among a million documents
- ◆ Naïvely, we would have to compute **pairwise Jaccard similarities** for every pair of signatures
 - 10^6 choose 2
 - Recall: for large n , $\binom{n}{2}$ is approximately $n^2/2$
 - $\approx 5 \cdot 10^{11}$ comparisons
 - At 10^5 secs/day and 10^6 comparisons/sec, it would take **6 days**.

Locality Sensitive Hashing Overview

- ◆ Hash items several times
 - In a way that **similar items** are more likely to be hashed to the same bucket than dissimilar items
- ◆ **Candidate Pair:** Any pair that hashes to the same bucket for **any** of the hashings
- ◆ Check only the candidate pairs for similarity
- ◆ **False positives** : **dissimilar pairs** that hash to the same bucket
- ◆ **False negatives:** **truly similar** pairs do **not** hash to the same bucket for at least one of the hash functions.

Recall: Minhashing Example



LSH: First Cut

- ◆ **Goal: Find documents with Jaccard similarity at least s** for some similarity threshold s (e.g. $s=0.8$)
- ◆ **LSH – General idea:** Use a function $f(x,y)$ that tells whether x and y are a **candidate pair**: a pair of elements whose similarity must be evaluated
- ◆ **For Min-Hash matrix:**
 - Hash columns of **signature matrix M** to many buckets
 - Each pair of documents that hashes into the same bucket is a **candidate pair**.

Candidates from Min-Hash

2	1	4	1
1	2	1	2
2	1	2	1


signature matrix **M**

- ◆ Pick a similarity threshold s ($0 < s < 1$)
- ◆ Columns x and y of M are a **candidate pair** if their signatures agree on at least fraction s of their rows:
 $M(i, x) = M(i, y)$ for at least frac. s values of I
 - We expect documents x and y to have the same (Jaccard) similarity as their signatures.


LSH for Min-Hash

- ◆ **Big idea:** Hash columns of signature matrix M several times
- ◆ Arrange that (only) **similar columns** are likely to **hash to the same bucket**, with high probability
- ◆ **Candidate pairs are those that hash to the same bucket.**


Partition M into b Bands



Partition M into Bands



- ◆ Divide matrix M into b bands of r rows
- ◆ For each band, hash its portion of each column to a hash table with k buckets
 - Make k as large as possible
 - Use a **separate bucket array for each band** so columns with the same vector in **different bands don't hash to same bucket**
- ◆ **Candidate column pairs** are those that **hash to the same bucket for ≥ 1 band**
- ◆ Tune b and r to catch most similar pairs, but few non-similar pairs.



Example of Bands

Assume the following case:

- ◆ Suppose 100,000 columns of M
 - Correspond to **signatures** for 100,000 documents
- ◆ Signatures of 100 integers (rows)
 - Correspond to 100 **hash functions** used in minhashing
- ◆ 4 bytes per integer
- ◆ Therefore, signatures take 40Mb
- ◆ Choose $b = 20$ bands of $r = 5$ rows of integers/band
- ◆ **Goal:** Find pairs of documents that are at least $s = 0.8$ or 80% similar.

Analysis of Banding Technique (Function)

- ◆ Use b bands of r rows each
- ◆ **Pair of documents have Jaccard similarity t**
 - Probability that minhash signatures for the documents agree in any one particular row of the signature matrix is t
- ◆ Columns C_1 and C_2 in signature matrix have similarity t
- ◆ Pick any band (r rows)
 - Prob. that all rows in band are equal = t^r
 - Prob. that not all r rows are equal (some row in band is unequal) = $1 - t^r$
- ◆ Prob. that no band has rows that are all equal = $(1 - t^r)^b$
- ◆ Prob. that at least 1 band has rows that are all equal (which is the probability of being a candidate pair) = $1 - (1 - t^r)^b$

C_1, C_2 are 80% Similar false negatives ?

2	1	4	1
1	2	1	2
2	1	2	1

- ◆ Find pairs of $\geq s=0.8$ similarity, set $b=20$, $r=5$
- ◆ Assume: $\text{sim}(C_1, C_2) = t = 0.8$
 - Since $\text{sim}(C_1, C_2) \geq s$, we want C_1, C_2 to be a **candidate pair**
 - We want them to hash to at **least 1 common bucket** (at least one band is identical)
- ◆ Probability C_1, C_2 identical in one particular band:
 $t^r = (0.8)^5 = 0.328$
- ◆ Probability C_1, C_2 are **not** similar in all of the 20 bands:
 $(1 - t^r)^b = (1 - 0.328)^{20} = 0.00035$
 - i.e., about .035% of the 80%-similar column pairs are **false negatives** (truly similar pairs that we miss)
 - We would find 99.965% pairs of truly similar documents.¹⁴

C_1, C_2 are 30% Similar false positives ?

2	1	4	1
1	2	1	2
2	1	2	1

- ◆ Find pairs of $\geq s=0.8$ similarity, set $b=20$, $r=5$
- ◆ Assume: $\text{sim}(C_1, C_2) = t = 0.3$
 - Since $\text{sim}(C_1, C_2) < s$ we want C_1, C_2 to hash to **NO common buckets** (all bands should be different)
 - Should NOT be a candidate pair!
- ◆ Probability C_1, C_2 identical in one particular band:
 $t^r = (0.3)^5 = 0.00243$
- ◆ Will identify C_1, C_2 as candidate pair if they are **identical in at least one band**
- ◆ Probability C_1, C_2 identical in at least 1 of 20 bands:
 $1 - (1 - t^r)^b = 1 - (1 - 0.00243)^{20} = 0.0474$
 - Approximately **4.74% pairs of docs with similarity 0.3** end up becoming **candidate pairs**
 - They are **false positives** (dissimilar documents that must be examined as candidate pairs but will have similarity below threshold s).

LSH Involves a Tradeoff

◆ Pick:

- The number of Min-Hashes (rows of M)
- The number of bands b , and
- The number of rows r per band

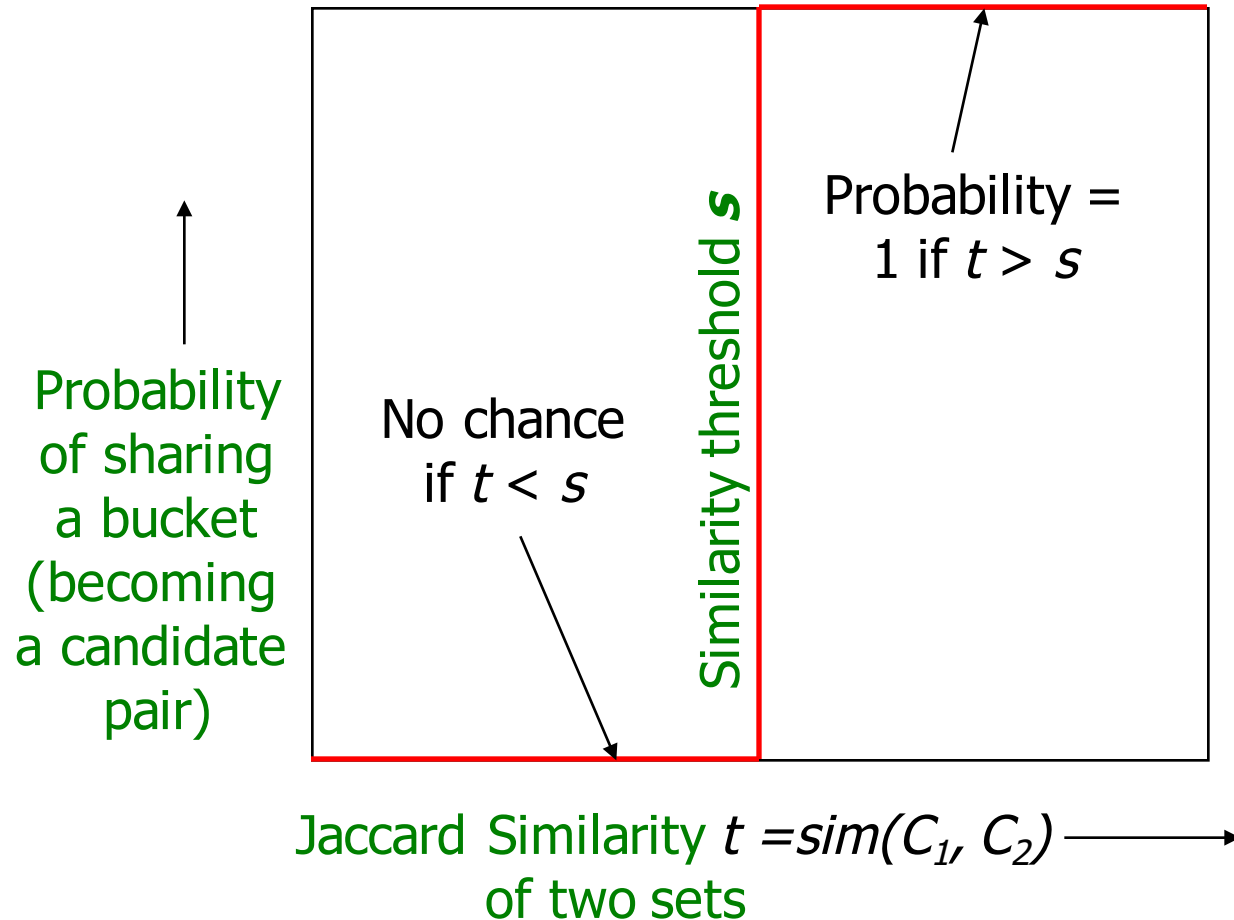
to balance false positives/negatives

- ◆ **Example:** If we had only 15 bands of 5 rows, the number of **false positives would go down**, but the number of **false negatives would go up**.

Example of Tradeoffs


- ◆ Previous example: 20 bands of 5 rows each
 - Probability of **false negatives** when C1, C2 are 80% similar:
0.00035
 - Probability of **false positives** when C1, C2 are 30% similar:
0.0474
- ◆ What if we use 15 bands of 5 rows each (smaller signature matrix)?
 - Probability of **false negatives higher** when C1, C2 are 80% similar:
 - $(1 - t^r)^b = (1 - 0.328)^{15} = 0.002573$
 - Probability of **false positives lower** when C1, C2 are 30% similar:
 - $1 - (1 - t^r)^b = 1 - (1 - 0.00243)^{15} = 0.0358.$

Analysis of LSH – What We Want




What 1 Band of 1 Row Gives You

- Compare two values in similarity matrix




What 1 Band of 1 Row Gives You

- Compare two values in similarity matrix



What b Bands of r Rows Gives You: $1 - (1 - t^r)^b$

- Form of an S-curve, regardless of values of b and r
- Threshold s is where rise of curve is steepest: approximately $(1/b)^{1/r}$



$R=5, b=20$, for $t=0.9$: $1 - (1 - t^r)^b = 0.99999$, for $t=0.1$: 0.0000199 ²¹

Example: $b = 20; r = 5$

- ◆ Similarity t of two columns
- ◆ Prob. that at least 1 band is identical (so a candidate pair):


t	$1-(1-t^r)^b$
.2	.006
.3	.047
.4	.186
.5	.470
.6	.802
.7	.975
.8	.9996

- Not an ideal step function
- Probability rises by more than 0.6 going from similarity $t = 0.4$ to $t = 0.6$
- Slope in middle > 3

Picking r and b : The S-curve

◆ Picking r and b to get the best S-curve

- 50 hash-functions ($r=5$, $b=10$)



Green area:

False Negative rate

Similar documents that are not identified as candidate pairs


Blue area:

False Positive rate

Dissimilar documents that are identified as candidate pairs

Picking b and r

- ◆ Threshold s defines how similar documents have to be for them to be regarded as a similar pair (e.g., $s = 0.8$)
- ◆ Length n for minhash signatures
- ◆ Pick number of bands b and number of rows r such that $br = n$ and threshold s is approximately $(1/b)^{1/r}$
- ◆ To avoid false negatives (green area):
 - Select b and r to produce a threshold lower than s
- ◆ To avoid false positives (blue area):
 - Select b and r to produce a higher threshold than s



Example
: $n=100$

b	r	$(1/b)^{1/r}$
50	2	0.1414
20	5	0.5493
10	10	0.7943
5	20	0.9227

Example

- ◆ $(1/b)^{1/r}$ represents the threshold of the S curve for function $1 - (1 - t^r)^b$, the probability of being a candidate pair
- ◆ If $s=0.6$ (similarity of documents to be a candidate pair) what values should you choose for b and r to reduce the number of **false negatives**?
- ◆ **To avoid false negatives:** Select b and r to produce a **threshold lower than s**
- ◆ To avoid **false positives**: Select b and r to produce a higher threshold than s
- ◆ Could choose $(b=20, r=5)$ or $(b=50, r=2)$: both give **threshold lower than s**
- ◆ **Better answer probably $b=20, r=5$**
- ◆ Because $b=50, r=20$ will have a higher rate of false positives: **TRADEOFFS**

Example
: $n=100$

b	r	$(1/b)^{1/r}$
50	2	0.1414
20	5	0.5493
10	10	0.7943
5	20	0.9227

LSH Summary

- ◆ Tune M , b , r to identify almost all candidate pairs with similar signatures, but eliminate most pairs that do not have similar signatures
- ◆ Then check in main memory that candidate pairs really do have similar signatures
- ◆ **Optional:** In another pass through data, check that the remaining candidate pairs really represent similar documents.

Summary: 3 Steps

- ◆ **Shingling:** Convert **documents to sets**
 - We used hashing to assign each shingle an ID
- ◆ **Min-Hashing:** Convert large **sets to short signatures**, while **preserving similarity**
 - We used **similarity preserving hashing** to generate signatures with property $\Pr[h_{\pi}(C_1) = h_{\pi}(C_2)] = \text{sim}(C_1, C_2)$
 - We used hashing to get around generating **random permutations**
- ◆ **Locality-Sensitive Hashing:** Focus on **pairs of signatures** likely to be from similar documents
 - We used hashing to find **candidate pairs** of similarity $\geq s$.

Combining the techniques (1)


1. Pick a value of k and **construct from each document the set of k -shingles**
 - Optionally **hash the k -shingles** to shorter bucket numbers
2. **Sort the document-shingle pairs** to order them by shingle
 - Which sets contain which elements (shingles)
3. Pick a length n for minhash signatures corresponding to n minhash functions and **compute the minhash signatures** for all the documents.

Combining the techniques (2)

4. **Choose threshold s** that defines how similar documents have to be for them to be regarded as a “similar pair”
 - **Pick number of bands b and number of rows r** such that $br = n$
 - **Adjust b and r to limit false positives or negatives**
5. **Construct candidate pairs with LSH technique**
6. **Examine candidate pair signatures** and determine whether fraction of components where they agree is at least s
7. **Optionally**, if signatures are sufficiently similar, **compare documents** to check they are truly similar.

Locality Sensitive Hashing

- ◆ *Or Near-neighbor search*
- ◆ Minhashing is one example of a **family of functions** (the minhash functions) **that can be combined** (by the banding technique) **to distinguish strongly between pairs at a low distance from pairs at a high distance**
- ◆ Steepness of the S-curve reflects how effectively we can **avoid false positives and false negatives** among the candidate pairs
- ◆ Section 3.6: more general theory of Locality Sensitive Functions



Families of Functions for LSH

- ◆ **Families of functions** (including minhash functions) that can serve to **produce candidate pairs efficiently**
 - Space of sets and Jaccard distance OR other space and/or distance measure
- ◆ **Three conditions for family of functions:**
 1. **More likely to make close pairs be candidate pairs than distant pairs**
 2. **Statistically independent**
 3. **Efficient** in two ways
 1. **Be able to identify candidate pairs in time much less than time to look at all pairs**
 2. **Combinable to build functions better at avoiding false positives and negatives** (e.g., banding technique takes single minhash functions, combines them to produce S-curve shape we want)

Locality-Sensitive Functions




Figure 3.9: Behavior of a (d_1, d_2, p_1, p_2) -sensitive function

LS Families of Hash Functions




Figure 3.9: Behavior of a (d_1, d_2, p_1, p_2) -sensitive function

- ◆ Suppose we have a **space S of points** with a **distance measure d**
- ◆ A **family H of hash functions** is said to be **(d_1, d_2, p_1, p_2) -sensitive** if for any x and y in S :
 1. If $d(x, y) \leq d_1$, then prob. over all h in H , that $h(x) = h(y)$ is at least p_1
 2. If $d(x, y) \geq d_2$, then prob. over all h in H , that $h(x) = h(y)$ is at most p_2
- ◆ Note: we say nothing about what happens when the distance between items is between d_1 and d_2
 - But can make d_1 and d_2 as close as we wish
 - Can drive p_1 and p_2 apart while keeping d_1 and d_2 fixed.

Locality Sensitive Hashing for Other Distance Measures

- ◆ We focused on **minhashing**, a locality sensitive hashing family that uses **Jaccard distance**
 - Based on sets representing documents and their Jaccard similarity
- ◆ Book covers LSH families for **other distance measures**:
 - **Euclidean distance**: based on the **locations of points in a *Euclidean space*** with some number of real-valued dimensions
 - **Cosine distance**: **angle between vectors** from the origin to the points in question
 - **Edit distance**: **number of inserts and deletes** to change one string into another
 - **Hamming Distance**: **number of positions** in which bit vectors differ


LSH and Shingling Application Examples

- ◆ Matching fingerprints
- ◆ Identifying similar news articles

LSH for Fingerprints

- ◆ Typical representation is not an image, but **set of locations in which minutiae are located**
 - **Place where something unusual happens**: two ridges merging or a ridge ending
- ◆ Place a grid over a fingerprint
 - Normalize for **size and orientation** so that identical prints will overlap
- ◆ Represent fingerprint by **set of grid points where minutiae are located**
 - Possibly, **treat minutiae near a grid boundary** as if also present in adjacent grid points.

Discretizing Minutiae



Place a minutia in several adjacent grid squares if it lies close to the border of the squares

Applying LSH to Fingerprints

- ◆ **Make a bit vector for each fingerprint's set of grid points with minutiae**
 - Similar to set representing a document: 1 if the shingle is in the document, 0 otherwise
- ◆ We could minhash the bit vectors to obtain signatures
 - But since there probably aren't too many grid points, we can work from the bit-vectors directly.

Matching Fingerprints with LSH:

Many-to-many problem

- ◆ **Many-to-many version of fingerprint matching:** take an entire database of fingerprints and identify if there are **any pairs** that represent the **same individual**
 - Analogous to finding similar documents among millions of documents
- ◆ **Define a locality-sensitive family of hash functions:**
 - Each function f in the family F is defined by 3 grid squares
 - Function f says “yes” for two fingerprints if both have minutiae in all three grid squares, otherwise, f says “no”
 - “Yes” means the two fingerprints are candidate pairs
- ◆ **Sort of “bucketization”**
 - Each set of three points creates one bucket
 - Function f sends fingerprints to its bucket that have minutae in all three grid points of f
- ◆ **Compare all fingerprints in each of the buckets.**

Matching Fingerprints with LSH: Many-to-One Problem

- ◆ **Many-to-one version** : A fingerprint has been found at a crime scene, and we want to **compare it with all fingerprints in a large database to see if there is a match**
- ◆ Could use many functions f from family F
- ◆ **Precompute their buckets of fingerprints to which they answer “yes” on the large database**
- ◆ **For a new fingerprint:**
 - **Determine which buckets it belongs to**
 - **Compare it with all fingerprints found in any of those buckets.**

Example 3.22

- ◆ 1024 functions chosen randomly from F
 - Each function f says “yes” for two fingerprints if both have minutiae in all three grid squares, otherwise, f says “no”
- ◆ Suppose **typical fingerprints have minutiae in 20% of the grid points**
- ◆ Suppose **fingerprints from the same finger agree in at least 80% of their points**
- ◆ *** Probability two random fingerprints each have 1 in all three points = $(0.2)^6 = .000064$**
 - 2 fingerprints, 3 points each, all independent events.

First image
has 1 in a
point

Example: Continued

Second image
of same finger
also has 1

- ◆ Probability **two fingerprints from the same finger each have 1's in three given points** =

$$((0.2)(0.8))^3 = .004096$$

(Analogy: t^r)

- ◆ Prob. for **at least one of 1024 sets of three points** = $1 - (1 - .004096)^{1024} = .985$

(Analogy:
 $1 - (1 - t^r)^b$)

- ◆ But for **random fingerprints** (prev. slide*):

$$1 - (1 - .000064)^{1024} = .063$$

**1.5% false
negatives**

**6.3% false
positives**

Choosing the number of functions from F

- ◆ **Want to use many functions from F, but not too many**
- ◆ **Want a good probability of matching fingerprints from the same finger while not having too many false positives**
- ◆ Previous example: only **1.5% chance we fail** to identify a print on the gun (**false negative**), but have to look at **6.3%** of entire database (due to **false positives**)
- ◆ **Increasing number of functions from F increases number of false positives**
 - Only a small benefit in reducing false negatives below 1.5%
- ◆ **Can use constructions/combinations of functions**
 - Several examples in the chapter.

Finding Same or Similar Same News Articles

- ◆ **Want to organize large repository of on-line news articles**
 - Group together web pages derived from same basic text
- ◆ **Scenario:** the same article, say from the Associated Press, appears on the Web site of many newspapers, but looks quite different
- ◆ **Each newspaper surrounds the text of the article with:**
 - Its own logo and text
 - Ads
 - Perhaps links to other articles
- ◆ **A newspaper may also “crop” the article (delete parts).**

Variation on Shingling


- ◆ **Looks like earlier problem:** find documents whose shingles have high Jaccard similarity
- ◆ **But: Shingling treats all parts of document equally**
- ◆ For this application, **we want to ignore parts of the documents** (ads, links to other articles, etc.)
- ◆ **There is a difference between text that appears in prose and text in ads or headlines/links**
 - Prose contains greater frequency of *stop words*
 - E.g., common words like “and” or “the”
 - Common to use list of several hundred most frequent words.

New Shingling Technique

- ◆ **News articles have a lot of stop words, while ads do not**
 - “Buy Sudzo” vs. “I recommend **that you** buy Sudzo **for your** laundry.”
- ◆ **Define a *shingle* to be a stop word plus the next two following words**
 - Shingles are: “I recommend **that**”
 - “**that you** buy”
 - “**you** buy Sudzo”
 - “**for your** laundry”
 - “**your** laundry <nextword>”
- ◆ **Then compare the similarity of the sets of shingles that represent each document**
 - Don’t use minhashing or LSH in this example.

Why it Works

- ◆ By requiring each shingle to have a stop word: **bias the mapping from documents to shingles** so it picked more shingles from the article than from the ads
- ◆ **Pages with the same article, but different ads,** have **higher Jaccard similarity** than those with the same ads, but different articles.



Locality Sensitive Hashing

Step 3: *Locality-Sensitive Hashing:*

Focus on **pairs of signatures** likely to be from similar documents